



Trelium Generalist / Ops Intern: Deep Research and Application Strategy

Executive conclusion

The strongest way to apply to this Trelium role is **not** to submit a better-looking résumé or a longer cover letter. It is to behave like the person the post is describing **before Trelium hires you**.

My recommendation is to spend roughly one day building a small, polished artifact called something like:

Trelium GTM Account & Workflow Intelligence Engine

A Codex/Claude Code-built tool that researches promotional-products companies, scores their fit for Trelium, identifies likely repetitive workflows, surfaces evidence, identifies the right buyer persona, and produces an account-specific GTM brief.

Then submit:

2–3 sentence email + résumé + LinkedIn + live demo/GitHub + one-page findings memo.

This is unusually well matched to the actual role. Trelium currently describes itself as an AI-agent platform that executes repetitive operational work across email, PDFs, spreadsheets, ERPs, CRMs and browser applications, with examples including order entry, quotation, invoice matching, customer updates, approval routing and reporting. It publicly targets mid-market and enterprise companies, while its recent product content shows a particularly strong vertical focus on promotional products, decorators, print shops, distributors and suppliers. ¹

Trelium is also very young. YC lists it as a Fall 2025 company founded in 2025, and its founders' public posts show a company that has been moving through customer discovery, conferences, launches and vertical specialization extremely quickly. ²

That means the applicant who demonstrates:

“Give me a vague growth/operations problem and I'll figure it out, build the solution, validate it, communicate what I learned, and hand you something usable”

is likely to be more compelling than the applicant who simply says:

“I am proactive, hardworking and good at Claude Code.”

That conclusion is an inference, but it is strongly supported by both the internship description you supplied and Trelium's broader hiring behavior. In another public hiring post, Trelium explicitly looked for people

who do not wait to be told what to do and people willing to try unusual growth ideas; Ritanshu's own hiring commentary also makes reliability and professional behavior especially salient. ³

My highest-confidence recommendation: build the GTM research engine, populate it with 15–30 real candidate accounts, manually validate the best five, record a 90-second demo, and let that be the centerpiece of your application.

Company, product, and apparent strategy

Trelium is a San Francisco-based, YC F25 applied-AI startup. YC identifies Abhimanyu Yadav as CEO and Ritanshu Dokania as a co-founder; YC's current jobs pages still label Ritanshu CTO, while Trelium's newer company materials have listed him as COO, so the public sources appear to be out of sync on his exact operating title. YC also reports a very small team, while Trelium's more recent company pages list additional team members, so it is safer to describe Trelium as a **small, early-stage team** rather than assert an exact employee count. ⁴

The company began with a narrower thesis. Ritanshu's September 2025 launch-era post described what was then Questom/Trelium as AI sales and support for commercial printing, targeting lost leads and manual customer follow-up. YC's launch material similarly described agents handling inbound B2B sales across phone, email, chat and SMS. The current product is considerably broader: Trelium now positions itself as an execution layer for repetitive operational processes across multiple business systems. The public record therefore suggests a **vertical customer problem → broader workflow platform evolution**, rather than the company simply remaining an inbound-sales agent. ⁵

There is still a strong vertical wedge. Trelium has become a PPAI-recognized solution provider and has explicitly said it is starting close to the promotional-products, print and custom-merchandise world before expanding into other industries characterized by heavy manual work. Recent content is highly specific to SAGE, PromoStandards, ShopWorks, Printavo, SanMar, S&S Activewear and related industry systems. ⁶

That vertical strategy matters enormously for your application. **Do not build a generic “AI lead generation agent.”** Trelium already has much deeper understanding of the segment than a generic prospecting tool would demonstrate.

What the product actually does

At a high level, a business describes a recurring workflow, Trelium connects the required systems, and an agent repeatedly performs the process. Its public use-case library currently lists 29 agents across operational categories, with examples such as business reporting, company-store-to-ERP sync, invoice follow-up, invoice matching, margin analysis, distributor order entry, order-status updates, PO creation and PO entry. ⁷

The product pattern is approximately:

Unstructured input → interpretation → business rules → system actions → exception handling → audit trail

Trelium says its agents can read email and documents, extract structured information, compare it against business rules, update systems of record, draft communications and escalate exceptions to humans. ⁸

Some concrete workflows illustrate the opportunity better:

Workflow	Typical starting point	Trelium action	Business outcome
Order entry	Email, PDF, form	Extract SKUs, quantities, decoration/shipping details and enter them into the order system	Less manual PO transcription. ⁹
Invoice matching	Supplier invoice	Compare invoice against order/PO and flag discrepancies	Finance reviews exceptions rather than every invoice. ¹⁰
Order status	Customer request	Look up status/tracking, update systems and prepare reply	Less manual status chasing. ⁸
Quoting	Email/request form	Retrieve product/inventory/pricing information, apply business rules and construct quote	Faster response to prospective customers. ¹¹
Reporting	ERP/accounting/CRM data	Aggregate operational information and send recurring report	Less spreadsheet/report assembly. ¹²
Supplier purchasing	Approved customer order	Validate item, inventory, pricing and supplier rules; create supplier order	Fewer purchasing errors. ¹³
Invoice follow-up	Accounting data	Detect overdue invoices and send contextual follow-ups	Less accounts-receivable chasing. ⁷
Proposal follow-up	CRM/SAGE/commonsku	Track open opportunities and trigger contextual follow-ups	Less leakage after quotation. ⁸

Trelium's SanMar/S&S material provides perhaps the clearest example of the underlying value. Supplier systems can already provide product information, inventory, prices, orders, invoices and shipment data. The difficult part for customers is connecting those facts into the rest of their business process—email → supplier → quoting → ERP/order system → accounting → customer communication. Trelium is selling the **workflow layer between those systems**, not trying to replace them. ¹³

That distinction should influence your application project as well.

Technology and product philosophy

Trelium emphasizes several architectural ideas:

Natural-language construction. Users can describe what should happen instead of manually programming every integration. ¹⁴

Deterministic execution where possible. Trelium's materials explicitly distinguish reasoning from predictable software actions, using AI where judgment is required and deterministic workflow steps where it is not. ¹⁵

Cross-system execution. Current integrations mentioned publicly include systems such as NetSuite, QuickBooks, Salesforce, SAP, Oracle, HubSpot, Printavo, ShopWorks, Gmail, Outlook, Slack and many others; current materials advertise more than 3,000 integrations plus browser-based automation where an API is unavailable. ¹⁶

Human review and observability. The platform emphasizes run history, human review gates and exception escalation rather than autonomous action with no oversight. ¹⁷

Organization-wide workflows rather than individual copilots. That separates Trelium's intended use from one-off conversational AI. ⁸

A clever application project should mirror that architecture. Your GTM engine should use Claude/Codex for ambiguous research and extraction, while keeping account scoring, schemas, validation and exports deterministic. You would thereby demonstrate that you understood **how Trelium itself thinks about AI systems**.

Business model

Trelium has a transparent SaaS/usage model. Its public pricing currently shows Starter around \$500 per month, Growth around \$1,000, Supreme around \$2,000, with discounted annual-equivalent pricing and custom Enterprise plans. Plans are allocated credits consumed by building and running agents and correspond approximately to increasing numbers of workflows/agents. ¹⁸

So the economic sale is not simply:

“Buy AI.”

It is closer to:

“Here is a recurring operational process consuming employee time every week. Automating it costs less than continuing to execute it manually.”

Trelium has increasingly published content around measuring workflow-level AI cost and ROI, which reinforces that outcome-oriented selling motion. ¹⁹

Current stage and strategy

Three things appear to define the current stage.

First, it is an early YC startup, founded in 2025 and still operating with a very small team. ²⁰

Second, it has moved substantially deeper into the branded-merchandise ecosystem. Trelium publicly lists customers such as Vapor Apparel, Vantage, HPG, Numo, iPromo, Brand Fuel, Icebox Cool Stuff and others, while recent integration content increasingly targets the technology used by distributors, suppliers, decorators and print shops. ²¹

Third, the founders continue to behave in a hands-on customer-development fashion. Abhimanyu publicly shared footage recalling a conference that the founders decided to attend at extremely short notice simply to meet potential customers, while Ritanshu's content repeatedly features industry events and operational examples. ²²

My interpretation: Trelium's current strategy looks like **win a specific operations-heavy vertical deeply, become embedded across many of its core workflows, then expand the workflow platform beyond that vertical**. The founders themselves have described promotional products as a starting point rather than necessarily the final market boundary. ²³

That is a much more useful mental model than “another AI-agent startup.”

Ideal customer profile and buying signals

There are two levels of ICP.

Broad ICP: mid-market and enterprise organizations with repetitive work spanning email, documents, spreadsheets and multiple systems of record. Trelium explicitly describes this audience on its site. ¹⁶

Current highest-confidence vertical ICP: North American promotional-products distributors, suppliers, decorators, print shops and related merchandise businesses that process substantial volumes of quotations, POs, customer orders, supplier interactions, invoices and status requests across fragmented software. Trelium's PPAI positioning, recent product content and integration strategy make this vertical concentration unusually clear. ⁶

Here is the practical ICP model I would use in your application project:

Dimension	Evidence	Working ICP hypothesis
Industry	Trelium explicitly targets distributors, suppliers, decorators and print shops, while also describing manufacturing and other operations-heavy sectors. ²⁴	Prioritize promotional products first.
Scale	Website says mid-market and enterprise. ¹⁶	Roughly established businesses with sufficient transaction volume to make recurring automation valuable.
Geography	Trelium is SF-based, participates heavily in U.S. industry associations/events, and its named promo customers are concentrated in North America. ²⁵	U.S./Canada is the safest initial target market.

Dimension	Evidence	Working ICP hypothesis
Systems	Public examples repeatedly mention SAGE, commonsku, ShopWorks, Printavo, NetSuite, QuickBooks, Gmail/Outlook and supplier APIs. ²⁶	Multiple disconnected operational systems is a strong fit signal.
Work	Quoting, PO intake, order entry, supplier purchasing, status updates, invoice matching, AR and reporting. ²⁷	Repetitive, high-frequency, rules-heavy work with some unstructured inputs.
Users	Trelium publicly describes use across operations, finance and GTM. ²⁸	Ops, finance/AP/AR, customer-service and sales-ops teams.
Buyer	Public testimonials include CFO, CEO and chief experience leadership. ¹⁶	COO/CEO at smaller firms; VP Ops/ CFO/technology/transformation leaders in larger companies. This is an inference.

The strongest pain pattern

The best Trelium customer probably does **not** merely have “manual work.”

Almost every company has manual work.

The stronger signal is:

A high-volume process begins with messy/unstructured information, crosses several software systems, has repeatable business rules, and still requires humans to handle exceptions.

Consider an incoming order:

```

Customer email
  ↓
PDF / Excel PO
  ↓
Extract item + quantity + artwork + shipping data
  ↓
Check supplier data
  ↓
Apply pricing / margin / inventory rules
  ↓
Create order in ShopWorks / ERP
  ↓
Attach files
  ↓

```

Update spreadsheet / CRM



Confirm with customer

Trelium publicly describes an active pattern very similar to this using incoming order emails, PDF/Excel attachments, SanMar information, ShopWorks and Google Sheets. ¹³

That gives you a much stronger research question than:

“Does this company use AI?”

Your engine should ask:

“Where does information cross system boundaries inside this company?”

Buying triggers worth researching

The following are **my inferred buying triggers**, derived from Trelium's demonstrated workflows and the industry's operational structure:

Trigger	Why it matters
Rapid company growth	Transaction volume may grow faster than operations headcount.
Acquisition/integration	Multiple processes and legacy systems create duplicate work.
New ERP/order platform	Migration often exposes process inefficiencies.
High customer-service volume	Repeated status/quote requests become automation candidates.
Large supplier ecosystem	More inventory/pricing/status checks across systems.
Hiring for operations/order-entry/AP roles	Possible signal that manual workload is increasing.
Multiple systems named publicly	Cross-system data movement is precisely Trelium's value proposition.
High quotation volume	Faster quotation and follow-up can connect automation directly to revenue.
Complex invoice reconciliation	Finance automation has measurable labor/accuracy value.
Margin pressure	Automation becomes easier to justify when operational cost reduction matters.

These are hypotheses rather than confirmed Trelium qualification criteria, but they follow directly from the workflows Trelium currently markets. ²⁷

Who you should identify

For your mini-project, do **not** spend time finding personal emails.

Identify roles:

For distributors: COO/VP Operations, President/CEO at smaller firms, VP Technology, Sales Operations, Finance/CFO.

For suppliers: VP Operations, Customer Experience/Service leadership, Finance/AP, CIO/VP Technology.

For decorators/print shops: owner/CEO, operations manager, production/customer experience leader.

Then tell Trelum *why that role* is likely to own the workflow.

That demonstrates considerably more GTM reasoning than merely presenting 30 names from Apollo.

GTM motion, founder signals, and what Trelum is actually hiring for

Trelum's GTM already appears much richer than pure cold outbound.

Founder-led customer acquisition

There is strong evidence of founder-led selling/customer discovery. The founders have publicly documented going to industry conferences to find customers, including deciding on attendance with very little lead time, and Ritanshu continues to post product examples directed specifically at print and promotional-product operators. ²²

That suggests the internship is probably not a neat corporate “operations internship.” You are likely to be supporting a founder-driven growth machine where an assignment could move from:

research this market

to

build the list

to

automate the research

to

launch the experiment

to

figure out why it failed

with little ceremony between stages.

That is an inference, but it aligns almost exactly with the hiring post you supplied.

Content and LinkedIn

The founders and GTM team publish frequent, vertical-specific educational and product content. Trelium's current blog includes articles on NetSuite automation, AI workflow ROI, SAGE, SanMar/S&S, Claude Cowork, Syncore and other operational topics. ²⁹

The content strategy is not generic "AI is changing business" thought leadership. Much of it targets **high-intent operational questions** such as integrating SAGE or supplier APIs, which also gives the site an obvious search/discovery function. Calling this a deliberate SEO program would be an inference, but the content structure is strongly compatible with SEO-led demand generation. ²⁶

Events and community

Public posts show participation in the print/promo ecosystem, including conferences and broader outreach journeys. Trelium has also used industry demos and event-related content as a GTM tool. ³⁰

This matters because a good intern experiment could amplify existing event activity instead of inventing an entirely disconnected acquisition channel.

Partnerships

Trelium has a public partnership motion aimed at implementation firms, integrators and related service providers that can design, deploy and support agents. ³¹

This suggests an important second route to market:

sell directly to companies + equip service/technology partners who already advise those companies.

That is a potentially high-leverage intern project.

Product-led growth

I found **weak evidence for classic self-serve PLG.**

Trelium publishes pricing, but its main conversion flow repeatedly directs prospects toward talking to a human, including a distinctive 30-minute demonstration where Trelium says it will automate one of the prospect's actual workflows using their actual tools rather than give a slide presentation. ¹⁶

Therefore, the current motion appears **sales-/demo-led**, even though the underlying product is configurable by users. That is an inference from the public funnel, not a statement from Trelium.

Referrals

There is anecdotal evidence of word-of-mouth. In a recent personal post, Ritanshu described an industry mentor putting in a good word for Trelium with a friend considering the product. That demonstrates relationship-driven referral behavior, although I found no evidence of a formalized public referral program.

32

What Ritanshu appears to value

The internship post you provided should remain the primary signal. But other public material reinforces several themes.

Self-direction. A separate Trelium hiring post emphasized candidates who do not wait to be told what to do and GTM-oriented hires willing to try unconventional experiments. 33

Speed and willingness to act under incomplete information. The founders themselves have publicly described making abrupt decisions to attend conferences simply to learn and meet customers. 34

Real customer understanding. Ritanshu's recent product content is highly concrete—actual print-shop quotation behavior rather than abstract AI language. 35

Reliability and professionalism. Ritanshu publicly complained about interview candidates who no-showed, behaved poorly after rejection or otherwise mishandled basic professional interactions. It would be unreasonable to infer an entire hiring rubric from one post, but it reinforces that “reliable” in your internship post is probably not filler language. 36

Persistence. In the original YC announcement, the founders described the archetype they wanted to represent as one rooted in stability and persistence, while also acknowledging that startup life occasionally requires intense bursts of effort. 37

What this means for your application

Do **not** lead with:

“I am hardworking and passionate about AI.”

Do not list 15 frameworks.

Do not write five paragraphs about why YC startups excite you.

Instead, demonstrate four things:

I noticed something → I investigated it → I built something useful → here is what I learned.

That is almost the perfect compressed demonstration of this job description.

Competitive landscape and Trelium's differentiation

Trelium operates in a crowded automation/agent category, so understanding what it is **not** is useful.

Alternative	What it is strong at	How Trelium appears positioned differently
Zapier	Very broad no-code automation platform with more than 9,000 app connections plus AI agents and workflows. ³⁸	Trelium emphasizes complete operational jobs and vertical workflows rather than primarily giving users a general-purpose automation toolkit. That distinction is partly Trelium's own positioning. ³⁹
n8n	Flexible workflow automation with AI agents, flow logic, human approvals, cloud and self-hosted deployment options. ⁴⁰	Trelium appears more targeted toward business operators describing a process and deploying an "employee," rather than technical teams assembling workflows. This is an inference from product UX/positioning.
UiPath	Mature enterprise automation/orchestration combining AI agents, deterministic automation, humans, governance and complex end-to-end processes. ⁴¹	Trelium's likely advantage is speed/simplicity/vertical knowledge for mid-market operational teams rather than UiPath's enterprise automation breadth. This is a positioning inference, not evidence that UiPath cannot serve Trelium's customers.
Salesforce Agentforce	Agents deeply integrated with Salesforce's CRM/data/application ecosystem and industry solutions. ⁴²	Trelium's story is system-agnostic execution across ERPs, email, vertical order systems, supplier infrastructure and browser applications rather than beginning with Salesforce as the platform center. ¹⁶
Custom integrations / scripts	Maximum custom control and can use existing vendor APIs directly.	Trelium argues the recurring problem is the engineering/maintenance burden needed to turn APIs into complete business processes. Its SanMar/S&S integration article describes exactly this gap. ¹³

An important caution: many of these companies are converging technologically. UiPath now explicitly combines agents with deterministic workflows; n8n supports human review; Zapier has AI agents; Salesforce is expanding beyond simple CRM automation. So **"Trelium has AI agents and competitors don't" is not a defensible differentiation thesis.** ⁴³

The more credible Trelium differentiation appears to be the combination of:

business-user workflow definition + cross-system execution + deterministic controls + human exception handling + aggressive verticalization around operational workflows.

That last part is especially important.

Zapier can connect thousands of applications. UiPath can automate extremely complex enterprises. n8n can build sophisticated agent workflows.

But Trelium can walk into a promotional-products company and say, conceptually:

“Show us how the PO comes in, how you check SanMar, what happens in ShopWorks, when accounting becomes involved, and where you respond to the customer.”

Its recent content demonstrates increasingly deep familiarity with exactly those workflows. 26

That vertical operational knowledge may be a more strategically meaningful wedge than any isolated technical feature.

And that again tells you what your application should demonstrate:

understand the customer's business, not merely the AI stack.

GTM experiments and prospect research

The following experiments are ranked using **my judgment**, not Trelium's internal roadmap. I prioritized experiments that one highly capable intern could initiate rapidly.

Rank	Experiment	Impact	Difficulty	Cost	First signal
1	Workflow Opportunity Audit outbound — research each target and send a personalized one-page hypothesis about a workflow Trelium could automate	Very high	Medium	<\$100	2–3 days
2	PPAI Account Intelligence Engine — automatically score leading distributors/suppliers using workflow + stack + growth signals	Very high	Medium	<\$100	1–2 days
3	Five-account video teardown test — record short personalized demonstrations explaining likely workflow opportunities	High	Medium	~\$0	1–2 days
4	SAGE/commonsku/ShopWorks signal outbound — identify companies mentioning relevant systems and target an integration-specific message	High	Medium	<\$100	2–4 days

Rank	Experiment	Impact	Difficulty	Cost	First signal
5	Partner activation sprint — identify implementation consultants, ERP/order-system specialists and promo-tech advisors for Trelium's partner program	High	Medium	~\$0	3–7 days
6	“Bring us your worst workflow” clinic — vertical webinar/office-hours version of Trelium's existing live-workflow demo	High	Medium	<\$200	1 week
7	Trigger-based account watchlist — monitor hiring, acquisitions, ERP changes and rapid growth among target accounts	Medium-high	Medium	<\$100	3–5 days
8	ROI-first outbound — estimate manual process cost and present an automation ROI hypothesis rather than a generic product pitch	Medium-high	Low	~\$0	2–3 days
9	Customer workflow expansion audit — systematically identify second/third automation opportunities within existing customers	Very high potential	Medium	~\$0	1 week
10	Customer proof/referral flywheel — turn completed workflows into tiny workflow-specific case studies and referral prompts	Medium-high	Medium	~\$0	1 week

The first two are best for **your application** because they require no internal Trelium access.

The ninth could ultimately have extremely high business value, but you cannot genuinely execute it before hiring because it requires customer and product data.

Why the workflow-audit experiment is particularly strong

Instead of emailing:

“Trelium helps automate your operations.”

A researched outbound asset could say:

“Your company appears to operate at significant promotional-products volume. Based on your business model and the systems/processes typical of companies in this segment, I'd investigate incoming PO → order entry → supplier validation as the first automation candidate. Here's the six-step workflow I'd validate on a discovery call.”

The key word is **hypothesis**.

Do not claim a company has a specific manual problem unless there is evidence.

That is also why your tool should visibly separate:

FACTS from **INFERENCES** from **QUESTIONS TO VALIDATE**.

That discipline alone would make the project substantially better.

Candidate prospect list

I screened Trelium's publicly displayed customer logos first and intentionally avoided those named accounts in the list below. Trelium's website currently displays customers including HPG, Vantage, Vapor Apparel, iPromo, Numo, Brand Fuel, Icebox Cool Stuff and numerous others. Because the website is unlikely to represent every customer or active opportunity, **all accounts below must be de-duplicated against Trelium's CRM before anyone contacts them.** ¹⁶

The approximate company-size figures below are **2025 promotional-products revenue reported/estimated in the 2026 PPAI 100**, not total corporate revenue in every case. The ranking is my own estimate of attractiveness for a research experiment, not PPAI's rank. ⁴⁴

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
1	Concord Marketing Solutions	Distributor; ~\$87M	PPAI reported very high recent growth, making operational scalability a plausible research angle. ⁴⁵	PO/order entry, reporting, invoice reconciliation	COO/ Operations: "Can operational throughput keep pace with growth without matching headcount growth?"
2	Stran Promotional Solutions	Distributor; ~\$116M	Large, rapidly growing PPAI-ranked distributor. ⁴⁵	Quote → order handoff, PO entry, reporting	Operations/ technology: "Where is growth creating manual handoff?"

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
3	Overture Promotions	Distributor; ~\$192M	Significant PPAI-ranked scale and strong reported growth. ⁴⁵	Orders, customer status, finance matching	COO: “Which high-volume customer workflow has become hardest to scale?”
4	iPROMOTEu	Distributor organization; ~\$372M	Large industry participant with strong reported growth. It is distinct from iPromo, which Trelum lists publicly as a customer. ⁴⁶	Order intake, distributed reporting, quote follow-up	Operations: “Can repetitive operational work be standardized across high transaction volume?”
5	American Solutions for Business	Distributor; ~\$335M	Large PPAI- ranked distributor with substantial transaction scale. ⁴⁵	PO entry, supplier communications, reporting	COO/tech leader: cross- system order- processing audit
6	BDA	Distributor; ~\$551M	Large, growing distributor where multi- system operational work is likely material; workflow is an inference. ⁴⁵	Orders, vendor status, invoices	Operations/IT: exception- based workflow execution
7	Geiger	Distributor; ~\$370M	Large established PPAI-ranked distributor. ⁴⁵	PO/order processing, supplier updates	Digital transformation/ operations: workflow inventory

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
8	HH Global	Distributor; ~\$457M promo revenue	Very large promotional- products operation with strong PPAI- reported growth. ⁴⁵	Reporting, order handoffs, invoice matching	Operations/ technology: target a defined cross-system workflow rather than broad “AI transformation”
9	Proforma	Distributor organization; ~\$668M	Very large promo business with high process volume. ⁴⁵	Order entry, quote follow-up, reporting	Operations/ technology: repetitive-work mapping
10	LeaderPromos	Distributor; ~\$59M	Mid-market scale may make Trelum's price/ implementation model easier to test than at a billion-dollar enterprise. Size from PPAI. ⁴⁵	Quote follow-up, order entry	President/ops: “Automate one painful process first.”
11	Nadel	Distributor; ~\$225M	Material PPAI- ranked volume suggests recurring operational processes worth researching. ⁴⁵	PO intake, invoice review, reporting	COO: identify highest-volume repetitive process
12	G&G Outfitters	Distributor; ~\$96M	Mid-market distributor with reported growth. ⁴⁵	Order entry/status	Operations: reduce time spent moving order information

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
13	HALO	Distributor; ~\$964M	One of the industry's largest distributors. ⁴⁵	High-volume exception handling	Enterprise operations/ transformation: narrowly scoped workflow POC
14	Staples Promotional Products	Distributor; ~\$930M estimated	Very large PPAI-listed operation. ⁴⁵	Reporting/order/ invoice workflow	Enterprise ops: identify workflow not already covered by internal automation
15	4imprint	Distributor; ~\$1.3B	Massive transaction scale, although likely substantial existing automation, so fit is less certain. ⁴⁵	Exception handling/ customer operations	Automation/ ops: search for gaps, not generic automation
16	Hit Promotional Products	Supplier; ~\$733M	Large PPAI-ranked supplier with substantial order volume. ⁴⁷	Distributor PO intake, status, invoice processing	VP Operations/ customer service: "Move from processing every request to managing exceptions."
17	Goldstar	Supplier; ~\$98M	PPAI reports significant growth; mid-market scale makes it an interesting automation hypothesis. ⁴⁷	PO intake/status/ invoices	Ops leader: scale growth without proportional administration

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
18	Showdown Displays	Supplier; ~\$150M	Large, growing supplier. ⁴⁷	PO entry, distributor status inquiries	Customer operations: remove repetitive order- status work
19	Ariel Premium Supply	Supplier; ~\$131M	Mid-sized, strongly growing PPAI- listed supplier. ⁴⁷	Incoming POs, order acknowledgments, invoice match	Operations: automate repetitive order administration
20	The Magnet Group	Supplier; ~\$130M	Significant supplier volume and reported growth. ⁴⁷	PO → ERP, status communication	Operations/ customer service: workflow- throughput audit
21	Hirsch	Supplier; ~\$59M	Mid-market supplier with strong reported growth. ⁴⁷	Orders/status/ finance	COO/ operations: quick one- workflow proof of concept
22	Ball Pro	Supplier; ~\$43M	Smaller PPAI- ranked supplier with exceptionally high reported growth, which is a strong trigger for operational research. ⁴⁷	PO entry and order-status updates	President/ops: "Where has the growth rate created manual bottlenecks?"
23	High Caliber Line	Supplier; ~\$50M	Fast-growing mid-market supplier according to PPAI. ⁴⁷	Incoming orders, customer service	Operations: growth-driven workflow automation
24	Koozie Group	Supplier; ~\$358M	Large PPAI supplier. ⁴⁷	PO processing, exception routing, finance	Digital ops: high-volume exception handling

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
25	Gemline	Supplier; ~\$143M	Meaningful supplier scale even though PPAI showed weaker recent growth. ⁴⁷	PO entry/status/ invoice workflow	Finance/ops: efficiency rather than growth angle
26	Logomark	Supplier; ~\$131M	Material industry supplier scale. ⁴⁷	Order entry and distributor communications	Customer ops: status/update workload
27	SnugZ USA	Supplier; ~\$92M	PPAI-ranked supplier; recent contraction makes an efficiency/cost angle more plausible than a scale-growth pitch. ⁴⁷	Finance/order administration	CFO/ operations: reduce repetitive operating cost
28	PCNA	Supplier; ~\$667M	Very large supplier, although enterprise sophistication means a targeted workflow is essential. ⁴⁷	Exceptions, supplier/ distributor data movement	Transformation/ operations: narrow POC

Priority	Candidate	Segment / approximate promo size	Fit hypothesis and evidence	Workflow hypothesis	Likely persona / outreach angle
29	S&S Activewear	Supplier; multi-billion-dollar	One of the industry's largest suppliers. Trelium already integrates S&S data into customer workflows, so S&S may be as much ecosystem infrastructure as prospect; treat this as low-confidence until checked internally. ⁴⁸	Internal cross-system operational workflow	Strategic/tech operations; do not use a generic supplier-data pitch
30	SanMar	Supplier; multi-billion-dollar	Industry-leading supplier and existing Trelium integration source; as with S&S, this could be partner/ecosystem rather than target. ⁴⁸	Internal exception-heavy ops	Strategic partnership/ops rather than ordinary outbound

Notice what this table intentionally does **not** do.

It does not claim:

"Concord manually enters every order."

We do not know that.

It says:

"Concord has characteristics that make order-entry/process scaling worth investigating."

That difference is critical.

The five accounts I would actually deep-dive for the application

Instead of trying to make 30 mediocre profiles, I would use the engine to score all 30 and then manually produce high-quality dossiers for:

Concord Marketing Solutions → Stran → Overture → Goldstar → Ball Pro.

Why those five?

They combine meaningful operational scale with unusually strong reported growth signals, while avoiding some of the enormous enterprises where existing automation programs may make a generic pitch less informative. The growth figures are from PPAI; my selection logic is an inference. 49

For each of those five, manually investigate:

Website → business model → products/services → operations language → job postings → public technology references → acquisitions/growth → likely workflow → buyer role → 3 discovery questions.

That deliverable would be far more impressive than dumping hundreds of scraped leads into a spreadsheet.

AI/Codex system you should build

Call it something straightforward:

Trelium GTM Intelligence

Do **not** make the project appear as though you are trying to replicate Trelium.

You are building an internal GTM research utility *for* Trelium.


Input

```
company_name  
company_domain  
optional_industry
```

Example:

Stran Promotional Solutions
stran.com
promotional products distributor

Research pipeline



What Claude/Codex should do

Use the model for things that involve interpretation:

- categorize company/segment;
- summarize business model;

- interpret public operational clues;
- infer likely workflows;
- produce discovery questions;
- explain account fit;
- generate possible outreach angles;
- distinguish fact from hypothesis.

What should remain deterministic

Do not let the LLM arbitrarily produce the final score.

Create an explicit rubric such as:

Signal	Weight
Core vertical fit	25
Operational/process complexity	20
Relevant software/ecosystem signals	15
Transaction/organization scale	15
Growth/buying trigger	15
Evidence quality	10
Total	100

Then implement rules.

Example:

```
promotional distributor/supplier/decorator: +25
strong evidence of multiple operating systems: +15
>$50M relevant revenue / material operation: +15
high-growth or acquisition trigger: +15
documented repetitive workflow clue: +20
multiple high-quality sources: +10
```

That architecture intentionally follows Trelum's own philosophy of using AI for uncertain interpretation and predictable software/rules for deterministic work. Trelum publicly promotes that separation between reasoning and deterministic execution. ¹⁵

Point that out in the README.

It will demonstrate that you did more than skim the homepage.

Required output schema

Every company should produce something like:

```
{
  "company": "Example Company",
  "segment": {
    "value": "promotional products distributor",
    "confidence": 0.94
  },
  "size": {
    "value": "$100M-$250M relevant revenue",
    "evidence": ["source"]
  },
  "fit_score": 84,
  "verified_facts": [
    {
      "claim": "...",
      "source": "...",
      "date": "..."
    }
  ],
  "stack_signals": [
    {
      "tool": "SAGE",
      "confidence": 0.72,
      "evidence": "..."
    }
  ],
  "buying_signals": [
    {
      "signal": "rapid growth",
      "evidence": "...",
      "confidence": 0.92
    }
  ],
  "workflow_hypotheses": [
    {
      "workflow": "PO -> order entry",
      "reason": "...",
      "confidence": "medium",
      "validation_questions": [
        "...",
        "...",
        "..."
      ]
    }
  ]
}
```

```

    ],
    "buyer_personas": [
      "VP Operations",
      "COO"
    ],
    "outreach_angle": "...",
    "research_gaps": [
      "No confirmed ERP found"
    ]
  }

```

The **research_gaps** field is important.

Most AI lead-research demos try to hide uncertainty.

You should surface it.

Source hierarchy

Your engine should prioritize roughly:

```

Company website
  ↓
Industry associations / filings / reputable databases
  ↓
Company job descriptions
  ↓
Executive/public professional content
  ↓
Reliable third-party reporting
  ↓
Search snippets / weaker evidence

```

Every fact should retain its source.

Never allow:

Inference → silently converted into fact

Instead:

FACT:
Company is a PPAI-ranked distributor.

INFERENCE:

Order-entry automation may be relevant because of its scale/industry.

VALIDATION QUESTION:

How are emailed purchase orders currently moved into the order-management system?

This is exactly the kind of research discipline that a good GTM generalist should display.

Interface

Do not waste half your day building a beautiful SaaS dashboard.

A clean CLI or minimal Next.js interface is enough.

A useful page could be:

```
| Research account  
| [ stran.com
```

```
] [Run] |
```

ICP FIT

87/100

Promotional distributor

HIGH CONFIDENCE

Operational complexity

HIGH

Growth trigger

HIGH

Relevant stack evidence

MEDIUM

TOP AUTOMATION OPPORTUNITY

PO / email → order-management workflow

Confidence: MEDIUM

Why:

- Large transaction-oriented distributor
- Rapid reported growth
- Workflow common in Trelum's current ICP

Evidence

[1] PPAI ...

[2] Company ...

[3] ...

VALIDATE BEFORE OUTREACH

1. How are incoming POs currently processed?
2. Which order-management system owns the record?
3. Where does manual re-entry still occur?

PERSONA

C00 / VP Operations

OUTREACH ANGLE

...

What the repository should contain

Keep it professional:

```
trelum-gtm-intelligence/  
|  
├─ README.md  
├─ src/  
|   ├─ research/  
|   ├─ extraction/  
|   ├─ scoring/  
|   ├─ workflows/  
|   └─ output/  
|  
├─ schemas/  
|   └─ account.ts  
|  
├─ examples/  
|   ├─ stran.json  
|   ├─ concord.json  
|   └─ overture.json  
|  
├─ data/  
|   └─ prospects.csv  
|  
└─ docs/  
    └─ findings.md
```

The README should be short:

Problem

Trelum's GTM team needs to determine which accounts have operational workflows worth automating.

What I built

A research pipeline that turns a company/domain into an evidence-backed account brief.

How it works

Public research → structured extraction → deterministic scoring → workflow hypothesis → GTM brief.

Why I built it this way

Mirroring Trelium's philosophy, the model handles ambiguous interpretation while deterministic code handles scoring, validation and output.

What I learned

Three concise observations from the prospect sample.

What I would do next

Test the best hypothesis against real discovery/outbound data.

That last section shows ownership beyond “code finished.”

The application I would submit

Your mini-project should be framed as an **experiment**, not a grand unsolicited consulting report.

The strongest title is probably:

I researched Trelium's ICP and built a Codex-powered account research system to find workflows worth automating.

Not:

“Strategic Growth Plan for Trelium.”

You are not inside the company. Pretending you know their GTM better than they do can backfire.

The finished package

By submission time, I would have exactly these assets:

Deliverable	What it demonstrates
Working research tool	Claude/Codex proficiency
GitHub repository	Execution + technical organization
30-account scored dataset	Market research + operational discipline
Five manually verified briefs	Quality control + GTM judgment

Deliverable	What it demonstrates
One-page findings memo	Communication
60–90 second demo	Ability to explain work
2–3 sentence email	Ability to follow instructions
Résumé + LinkedIn	Credentials/context

What to put in the one-page findings memo

Do not create a 20-page consulting deck.

Use:

Trelium GTM Research — What I Found

Question

Can public signals identify promotional-products companies with potentially strong Trelium workflow opportunities?

Then:

Method

Researched 30 PPAI suppliers/distributors, extracted scale/growth/technology/process evidence, scored accounts using deterministic criteria and manually reviewed the top five.

Then three findings, such as:

Finding A — Growth appears to be a useful prioritization signal.

Rapidly growing distributors may be especially interesting because transaction volume can increase faster than operational processes mature.

Finding B — Workflow specificity is more useful than generic AI personalization.

“Incoming PO → order entry” is a more actionable hypothesis than “you could benefit from AI.”

Finding C — Evidence quality is the bottleneck.

Public information rarely establishes whether a process is actually manual, so the system should generate validation questions instead of treating workflow hypotheses as facts.

Then:

Next experiment

Manually test five evidence-backed workflow hypotheses and compare response rates against generic account-personalized outbound.

That is a thoughtful GTM artifact.

The demo

Keep it around 60–90 seconds.

A good structure:

0–10 seconds

“I saw Trelium was hiring someone who can independently execute GTM experiments, so instead of only sending a résumé, I built one.”

10–25 seconds

Show input:

stran.com

25–50 seconds

Show output:

fit score → sources → buying signal → workflow hypothesis → buyer role.

50–65 seconds

Show the deterministic scoring/code.

“I use the model for interpreting messy company information, but keep scoring deterministic —similar to Trelium's own reasoning-versus-execution philosophy.”

65–80 seconds

Show prospect dataset.

“I ran 30 PPAI accounts and manually reviewed the five highest-priority opportunities.”

80–90 seconds

Finish:

“The main thing I learned is that workflow-specific research looks much more actionable than generic account personalization. Repo and findings are attached.”

Stop.

No eight-minute product walkthrough.

The email

Because Ritanshu specifically requested a **2-3 line note**, obey that constraint.

My preferred version:

Subject: Generalist / Ops Intern — [Your Name]

Hi Ritanshu — I saw you wanted someone proactive with Codex/Claude Code, so instead of only sending a résumé I built a small Trelium GTM experiment: it researches promo companies, scores ICP fit, and surfaces evidence-backed workflows worth investigating. I ran it across 30 candidate accounts and manually reviewed the top five: **[demo]** | **[GitHub]** | **[LinkedIn]** | **[resume]**

That's enough.

Another version, slightly bolder:

Hi Ritanshu — I used Codex to build a GTM research system for Trelium that turns a company into an evidence-backed ICP score, workflow hypothesis, buyer persona and outreach angle. I tested it on 30 promotional-products accounts and wrote up what I learned: **[demo/project]**. **LinkedIn:** **[link]** | **Resume:** **[link]**

Do **not** add:

“I hope this email finds you well...”

Do not write four paragraphs afterward.

Following the requested format itself signals communication discipline.

Your LinkedIn comment

After the application is sent:

Built a small Trelium-specific GTM experiment with Codex after seeing this — sent it over with my application. Really enjoyed digging into the workflow/ops side of the space!

That creates curiosity without publicly overselling yourself.

Your prioritized 24-hour execution plan

Assuming you can dedicate one focused day, I would allocate the effort like this:

Time	Work	Output
Hour 0–1	Read Trelum's homepage, use cases, current blogs, PPAI material and competitor positioning	1-page internal notes
Hour 1–2	Define ICP schema and scoring rubric before coding	JSON schema + score rules
Hour 2–5	Build research/extraction pipeline using Codex/Claude Code	Functional CLI/app
Hour 5–6	Add citations, confidence labels and fact/inference separation	Evidence layer
Hour 6–7	Add deterministic ICP scoring	Reproducible score
Hour 7–9	Run initial prospect set and debug hallucinations/data issues	First 30 accounts
Hour 9–11	Manually research the top five accounts	Five trustworthy briefs
Hour 11–12	Improve outreach/workflow hypothesis generation based on manual findings	Better outputs
Hour 12–13	Clean repository and write concise README	Public-ready GitHub
Hour 13–14	Write one-page findings memo	GTM artifact
Hour 14–15	Create minimal visual interface if needed	Demo-friendly UI
Hour 15–16	Test from clean install/input and fix broken paths/errors	Reliable build
Hour 16–17	Record 60–90 second demo	Loom/video
Hour 17–18	Update résumé/LinkedIn so AI + execution projects are immediately visible	Application assets
Hour 18	Send the 2–3 line email and then post the short LinkedIn comment	Application submitted

The exact elapsed hours matter less than the sequencing: **research enough to understand the problem, define the experiment, build, run it against real data, manually validate it, communicate what you learned, then submit.**

What I would not do

I would **not** spend the day building another generic RAG chatbot.

I would not create fake Trelium customer data.

I would not scrape hundreds of personal emails.

I would not send outbound messages to Trelium's potential customers without permission.

I would not imply that your hypothesized operational problems are known facts.

I would not produce AI-generated "personalization" such as:

"Saw your company is doing amazing things in promotional products..."

I would not spend six hours making a beautiful landing page while the underlying research is shallow.

And I would definitely not claim:

"Excellent with Claude Code/Codex"

without showing something you actually built with it.

The project itself should make that sentence unnecessary.

The application thesis should be:

You asked for someone who can take an ambiguous GTM/operations problem and execute without constant direction. I treated the application as my first assignment.

That is the strongest single way I see to turn this from an ordinary internship application into evidence that you already work the way Trelium says it needs.

1 8 16 17 18 21 25 31 Proposal Follow-Up Automation for Open Opportunities

https://www.trelium.com/?utm_source=chatgpt.com

2 4 20 Trelium: The Agent to Build Your Agents | Y Combinator

<https://www.ycombinator.com/companies/trelium>

3 33 We're building something big at Trelium (YC F25), and we need people who don't wait to be told what to do. Three roles we're hiring for right now: 1. Backend SWE - You geek out on systems. You have... | Ritanshu Dokania | 10 comments

https://www.linkedin.com/posts/ritanshudokania_were-building-something-big-at-trelium-activity-7454548676980645889-nQjb

5 37 Ritanshu Dokania's Post - LinkedIn

https://www.linkedin.com/posts/ritanshudokania_all-ive-ever-wanted-to-do-was-to-run-my-activity-7378818988954304512-RkmM?utm_source=chatgpt.com

6 24 Trelium Is Now a PPAI Recognized Solution Provider | Trelium Blog

https://www.trelium.com/blog/trelium-ppai-recognized-solution-provider?utm_source=chatgpt.com

7 27 AI Use Cases for Enterprises - Trelium

https://www.trelium.com/use-cases?utm_source=chatgpt.com

9 Order Entry AI Agent for Distributors | Trelium

<https://www.trelium.com/use-cases/order-entry-ai-agent-for-distributors>

10 Invoice Vouching Agent | Trelium

<https://www.trelium.com/use-cases/po-invoice-matching-agent>

11 13 26 How to Integrate SanMar and S&S APIs | Trelium

https://www.trelium.com/blog/integrate-sanmar-ss-activewear-api-trelium?utm_source=chatgpt.com

12 Business Report Agent | Trelium

<https://www.trelium.com/use-cases/business-report-agent>

14 15 AI Agent Platform Features for Enterprises | Trelium

<https://www.trelium.com/features>

19 29 Not Boring Blogs by Trelium

https://www.trelium.com/blog?utm_source=chatgpt.com

22 30 34 Ritanshu Dokania is out doing the Trelium (YC F25 ... - LinkedIn

https://www.linkedin.com/posts/abh1y1_ritanshu-dokania-is-out-doing-the-trelium-activity-7459980832695398400-M5OS?utm_source=chatgpt.com

23 About Trelium: AI Agents for Repetitive Work

<https://www.trelium.com/about>

28 Media Kit & Press Resources | Trelium

https://www.trelium.com/media?utm_source=chatgpt.com

32 Ritanshu Dokania's Post - LinkedIn

https://www.linkedin.com/posts/ritanshudokania_i-like-to-think-i-have-become-fairly-good-activity-7497687319659864065-KSUR?utm_source=chatgpt.com

35 Trelium AI Agents for 24/7 Quote Requests | Ritanshu Dokania ...

https://www.linkedin.com/posts/ritanshudokania_printshop-screenprinting-decoratedapparel-activity-7460309246216237057-LKqB?utm_source=chatgpt.com

36 Ritanshu Dokania's Post - LinkedIn

https://www.linkedin.com/posts/ritanshudokania_running-trelium-yc-f25-has-taught-me-a-activity-7455273460907528192--SN7?utm_source=chatgpt.com

38 43 How to build safe and trustworthy AI agents with Zapier

https://zapier.com/?utm_source=chatgpt.com

39 Trelium vs Zapier for Operations Automation

<https://www.trelium.com/compare/zapier>

40 Human-in-the-loop for tools | Build

https://docs.n8n.io/?utm_source=chatgpt.com

41 About agents - UiPath Documentation

https://docs.uipath.com/agents/automation-cloud/latest/user-guide/about-agents?utm_source=chatgpt.com

42 Agentforce Guide to Achieving Reliable Agent Behavior

https://www.salesforce.com/?utm_source=chatgpt.com

44 45 46 49 PPAI 100 - PPAI - Promotional Products Association International

<https://www.ppai.org/media/ppai-100/>

47 PPAI 100 Suppliers - PPAI - Promotional Products Association International

<https://www.ppai.org/media/ppai-100/suppliers/>

48 S&S Activewear: Into The Big Room - PPAI

https://www.ppai.org/media-hub/ss-activewear-into-the-big-room/?utm_source=chatgpt.com